

In [2]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import tensorflow as tf
```

In [3]: `df = pd.read_csv('Churn_Modelling.csv')`
`df.head()`


Out[3]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Ba
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	838
2	3	15619304	Onio	502	France	Female	42	8	1596
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	1255

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   RowNumber             10000 non-null  int64  
1   CustomerId            10000 non-null  int64  
2   Surname               10000 non-null  object  
3   CreditScore           10000 non-null  int64  
4   Geography             10000 non-null  object  
5   Gender                10000 non-null  object  
6   Age                   10000 non-null  int64  
7   Tenure                10000 non-null  int64  
8   Balance               10000 non-null  float64 
9   NumOfProducts        10000 non-null  int64  
10  HasCrCard             10000 non-null  int64  
11  IsActiveMember       10000 non-null  int64  
12  EstimatedSalary       10000 non-null  float64 
13  Exited                10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [5]: plt.xlabel('Exited')  
plt.ylabel('Count')  
df['Exited'].value_counts().plot.bar()  
plt.show()
```



```
In [6]: df['Geography'].value_counts()
```

```
Out[6]: France      5014  
Germany    2509  
Spain      2477  
Name: Geography, dtype: int64
```

```
In [7]: df = pd.concat([df, pd.get_dummies(df['Geography'], prefix='Geo')], axis=1)
```

```
In [8]: df = pd.concat([df, pd.get_dummies(df['Gender'])], axis=1)
```

In [9]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname                10000 non-null  object
3   CreditScore            10000 non-null  int64
4   Geography              10000 non-null  object
5   Gender                 10000 non-null  object
6   Age                    10000 non-null  int64
7   Tenure                 10000 non-null  int64
8   Balance                10000 non-null  float64
9   NumOfProducts          10000 non-null  int64
10  HasCrCard              10000 non-null  int64
11  IsActiveMember         10000 non-null  int64
12  EstimatedSalary        10000 non-null  float64
13  Exited                  10000 non-null  int64
14  Geo_France              10000 non-null  uint8
15  Geo_Germany             10000 non-null  uint8
16  Geo_Spain               10000 non-null  uint8
17  Female                  10000 non-null  uint8
18  Male                    10000 non-null  uint8
dtypes: float64(2), int64(9), object(3), uint8(5)
memory usage: 1.1+ MB
```

In [10]: df.drop(columns=['RowNumber', 'CustomerId', 'Surname', 'Geography', 'Gender'],

In [11]: df.head()

Out[11]:

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Estin
0	619	42	2	0.00	1	1	1	
1	608	41	1	83807.86	1	0	1	
2	502	42	8	159660.80	3	1	0	
3	699	39	1	0.00	2	0	0	
4	850	43	2	125510.82	1	1	1	

Splitting Data

In [12]: y = df['Exited'].values
x = df.loc[:,df.columns != 'Exited'].values

In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,random_state=20,test_

Scaling Data

```
In [14]: from sklearn.preprocessing import StandardScaler
std_x = StandardScaler()
x_train = std_x.fit_transform(x_train)
x_test = std_x.transform(x_test)
```

```
In [15]: x_train.shape
```

```
Out[15]: (7500, 13)
```

Tensorflow Model - Neural Network Classifier

```
In [16]: import tensorflow as tf
from tensorflow.keras.layers import Dense, Conv1D, Flatten
from tensorflow.keras.models import Sequential, Model
```

```
In [17]: model=Sequential()
model.add(Flatten(input_shape=(13,)))
model.add(Dense(100,activation='relu'))
model.add(Dense(1,activation='sigmoid'))
```

```
In [18]: model.compile(optimizer='adam',metrics=['accuracy'],loss='BinaryCrossentro
```

```
In [19]: model.fit(x_train,y_train,batch_size=64,validation_split=0.1,epochs=100)
```

```
Epoch 1/100
106/106 [=====] - 1s 3ms/step - loss: 0.5127
- accuracy: 0.7621 - val_loss: 0.4269 - val_accuracy: 0.8240
Epoch 2/100
106/106 [=====] - 0s 2ms/step - loss: 0.4306
- accuracy: 0.8119 - val_loss: 0.4042 - val_accuracy: 0.8280
Epoch 3/100
106/106 [=====] - 0s 2ms/step - loss: 0.4129
- accuracy: 0.8204 - val_loss: 0.3825 - val_accuracy: 0.8400
Epoch 4/100
106/106 [=====] - 0s 2ms/step - loss: 0.3970
- accuracy: 0.8323 - val_loss: 0.3681 - val_accuracy: 0.8560
Epoch 5/100
106/106 [=====] - 0s 2ms/step - loss: 0.3826
- accuracy: 0.8413 - val_loss: 0.3493 - val_accuracy: 0.8653
Epoch 6/100
106/106 [=====] - 0s 2ms/step - loss: 0.3712
- accuracy: 0.8428 - val_loss: 0.3389 - val_accuracy: 0.8640
Epoch 7/100
106/106 [=====] - 0s 2ms/step - loss: 0.3600
- accuracy: 0.8440 - val_loss: 0.3280 - val_accuracy: 0.8640
```

```
In [20]: pred = model.predict(x_test)
```

```
79/79 [=====] - 0s 3ms/step
```

```
In [21]: y_pred = []  
        for val in pred:  
            if val > 0.5:  
                y_pred.append(1)  
            else:  
                y_pred.append(0)
```


```
In [22]: from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay
```

```
In [23]: accuracy_score(y_test, y_pred)
```

Out[23]: 0.8564

```
In [24]: cm = confusion_matrix(y_test, y_pred)  
        display = ConfusionMatrixDisplay(cm)  
        display.plot()
```

Out[24]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2a6b0bb7070>



```
In [25]: from sklearn.neural_network import MLPClassifier
```

```
In [26]: nn_classifier = MLPClassifier(hidden_layer_sizes=(100),activation='logistic')
nn_classifier.fit(x_train,y_train)
```

```
c:\Users\hp\anaconda3\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:691: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (300) reached and the optimization hasn't converged yet.
  warnings.warn(
```

```
Out[26]: MLPClassifier(activation='logistic', hidden_layer_sizes=100, max_iter=300)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [27]: y_pred2 = nn_classifier.predict(x_test)
```

```
In [28]: accuracy_score(y_pred=y_pred2,y_true=y_test)
```

```
Out[28]: 0.8648
```

```
In [29]: nn_classifier.score(x_test,y_test)
```

```
Out[29]: 0.8648
```

```
In [ ]:
```